

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
import statsmodels.api as sm
from statsmodels.graphics.regressionplots import influence_plot
```

In [2]:


```
# import dataset  
toyo=pd.read_csv('ToyotaCorolla.csv',encoding='latin1')  
toyo
```

Out[2]:

	Id	Model	Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met_
0	1	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	13500	23	10	2002	46986	Diesel	90	
1	2	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	13750	23	10	2002	72937	Diesel	90	
2	3	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	13950	24	9	2002	41711	Diesel	90	
3	4	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	14950	26	7	2002	48000	Diesel	90	
4	5	TOYOTA Corolla 2.0 D4D HATCHB SOL 2/3- Doors	13750	30	3	2002	38500	Diesel	90	
...
1431	1438	TOYOTA Corolla 1.3 16V HATCHB G6 2/3- Doors	7500	69	12	1998	20544	Petrol	86	
1432	1439	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	10845	72	9	1998	19000	Petrol	86	
1433	1440	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	8500	71	10	1998	17016	Petrol	86	

	Id	Model	Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met_
1434	1441	TOYOTA								
		Corolla								
		1.3 16V								
		HATCHB	7250	70	11	1998	16916	Petrol	86	
		LINEA								
		TERRA								
		2/3-...								
1435	1442	TOYOTA								
		Corolla								
		1.6 LB								
		LINEA	6950	76	5	1998	1	Petrol	110	
		TERRA								
		4/5-								
		Doors								

1436 rows × 38 columns



In [3]:

```
toyo.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 38 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                     1436 non-null   int64
1   Model                  1436 non-null   object
2   Price                  1436 non-null   int64
3   Age_08_04              1436 non-null   int64
4   Mfg_Month              1436 non-null   int64
5   Mfg_Year                1436 non-null   int64
6   KM                     1436 non-null   int64
7   Fuel_Type              1436 non-null   object
8   HP                     1436 non-null   int64
9   Met_Color              1436 non-null   int64
10  Color                  1436 non-null   object
11  Automatic              1436 non-null   int64
12  cc                     1436 non-null   int64
13  Doors                  1436 non-null   int64
14  Cylinders              1436 non-null   int64
15  Gears                  1436 non-null   int64
16  Quarterly_Tax          1436 non-null   int64
17  Weight                 1436 non-null   int64
18  Mfr_Guarantee           1436 non-null   int64
19  BOVAG_Guarantee         1436 non-null   int64
20  Guarantee_Period       1436 non-null   int64
21  ABS                    1436 non-null   int64
22  Airbag_1               1436 non-null   int64
23  Airbag_2               1436 non-null   int64
24  Airco                  1436 non-null   int64
25  Automatic_airco        1436 non-null   int64
26  Boardcomputer          1436 non-null   int64
27  CD_Player              1436 non-null   int64
28  Central_Lock           1436 non-null   int64
29  Powered_Windows        1436 non-null   int64
30  Power_Steering         1436 non-null   int64
31  Radio                  1436 non-null   int64
32  Mistlamps              1436 non-null   int64
33  Sport_Model            1436 non-null   int64
34  Backseat_Divider       1436 non-null   int64
35  Metallic_Rim           1436 non-null   int64
36  Radio_cassette         1436 non-null   int64
37  Tow_Bar                1436 non-null   int64
dtypes: int64(35), object(3)
memory usage: 426.4+ KB
```

In [4]:

```
toyo2=pd.concat([toyo.iloc[:,2:4],toyo.iloc[:,6:7],toyo.iloc[:,8:9],toyo.iloc[:,12:14],
toyo.iloc[:,15:18]],axis=1)
toyo2
```

Out[4]:

	Price	Age_08_04	KM	HP	cc	Doors	Gears	Quarterly_Tax	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1431	7500	69	20544	86	1300	3	5	69	1025
1432	10845	72	19000	86	1300	3	5	69	1015
1433	8500	71	17016	86	1300	3	5	69	1015
1434	7250	70	16916	86	1300	3	5	69	1015
1435	6950	76	1	110	1600	5	5	19	1114

1436 rows × 9 columns

In [5]:

```
toyo3=toyo2.rename({'Age_08_04':'Age','cc':'CC','Quarterly_Tax':'QT'},axis=1)
toyo3
```

Out[5]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1431	7500	69	20544	86	1300	3	5	69	1025
1432	10845	72	19000	86	1300	3	5	69	1015
1433	8500	71	17016	86	1300	3	5	69	1015
1434	7250	70	16916	86	1300	3	5	69	1015
1435	6950	76	1	110	1600	5	5	19	1114

1436 rows × 9 columns

In [6]:

```
toyo3[toyo3.duplicated()]
```

Out[6]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
113	24950	8	13253	116	2000	5	5	234	1320

In [7]:

```
toyo4=toyo3.drop_duplicates().reset_index(drop=True)  
toyo4
```

Out[7]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1430	7500	69	20544	86	1300	3	5	69	1025
1431	10845	72	19000	86	1300	3	5	69	1015
1432	8500	71	17016	86	1300	3	5	69	1015
1433	7250	70	16916	86	1300	3	5	69	1015
1434	6950	76	1	110	1600	5	5	19	1114

1435 rows × 9 columns

In [8]:

```
toyo4.describe()
```

Out[8]:

	Price	Age	KM	HP	CC	Doors
count	1435.000000	1435.000000	1435.000000	1435.000000	1435.000000	1435.000000
mean	10720.915679	55.980488	68571.782578	101.491986	1576.560976	4.032753
std	3608.732978	18.563312	37491.094553	14.981408	424.387533	0.952667
min	4350.000000	1.000000	1.000000	69.000000	1300.000000	2.000000
25%	8450.000000	44.000000	43000.000000	90.000000	1400.000000	3.000000
50%	9900.000000	61.000000	63451.000000	110.000000	1600.000000	4.000000
75%	11950.000000	70.000000	87041.500000	110.000000	1600.000000	5.000000
max	32500.000000	80.000000	243000.000000	192.000000	16000.000000	5.000000

In [9]:

```
# correlation analysis  
toyo4.corr()
```

Out[9]:


	Price	Age	KM	HP	CC	Doors	Gears	QT
Price	1.000000	-0.876273	-0.569420	0.314134	0.124375	0.183604	0.063831	0.211508
Age	-0.876273	1.000000	0.504575	-0.155293	-0.096549	-0.146929	-0.005629	-0.193319
KM	-0.569420	0.504575	1.000000	-0.332904	0.103822	-0.035193	0.014890	0.283312
HP	0.314134	-0.155293	-0.332904	1.000000	0.035207	0.091803	0.209642	-0.302287
CC	0.124375	-0.096549	0.103822	0.035207	1.000000	0.079254	0.014732	0.305982
Doors	0.183604	-0.146929	-0.035193	0.091803	0.079254	1.000000	-0.160101	0.107353
Gears	0.063831	-0.005629	0.014890	0.209642	0.014732	-0.160101	1.000000	-0.005125
QT	0.211508	-0.193319	0.283312	-0.302287	0.305982	0.107353	-0.005125	1.000000
Weight	0.575869	-0.466484	-0.023969	0.087143	0.335077	0.301734	0.021238	0.621988

In [10]:

```
sns.set_style(style='darkgrid')  
sns.pairplot(toyo4)
```

Out[10]:

<seaborn.axisgrid.PairGrid at 0x25ba0980760>



In [11]:

```
# model building
model=smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=toyo4).fit()
```

In [12]:

```
# model testing
# Finding Coefficient parameters
model.params
```

Out[12]:

```
Intercept    -5472.540368
Age           -121.713891
KM            -0.020737
HP            31.584612
CC            -0.118558
Doors         -0.920189
Gears         597.715894
QT            3.858805
Weight        16.855470
dtype: float64
```

In [13]:

```
# Finding tvalues and pvalues
model.tvalues , np.round(model.pvalues,5)
```

Out[13]:

```
(Intercept    -3.875273
Age           -46.551876
KM            -16.552424
HP            11.209719
CC            -1.316436
Doors         -0.023012
Gears         3.034563
QT            2.944198
Weight        15.760663
dtype: float64,
Intercept      0.00011
Age            0.00000
KM             0.00000
HP             0.00000
CC             0.18824
Doors          0.98164
Gears          0.00245
QT             0.00329
Weight         0.00000
dtype: float64)
```

In [14]:

```
# Finding rsquared values
model.rsquared , model.rsquared_adj # Model accuracy is 86.17%
```

Out[14]:

```
(0.8625200256947, 0.8617487495415146)
```

In [15]:

```
# Build SLR and MLR models for insignificant variables 'CC' and 'Doors'
# Also find their tvalues and pvalues
slr_c=smf.ols('Price~CC',data=toyo4).fit()
slr_c.tvalues , slr_c.pvalues # CC has significant pvalue
```

Out[15]:

```
(Intercept    24.879592
CC            4.745039
dtype: float64,
Intercept    7.236022e-114
CC           2.292856e-06
dtype: float64)
```

In [16]:

```
slr_d=smf.ols('Price~Doors',data=toyo4).fit()
slr_d.tvalues , slr_d.pvalues
```

Out[16]:

```
(Intercept    19.421546
Doors         7.070520
dtype: float64,
Intercept    8.976407e-75
Doors        2.404166e-12
dtype: float64)
```

In [17]:

```
mlr_cd=smf.ols('Price~CC+Doors',data=toyo4).fit()
mlr_cd.tvalues , mlr_cd.pvalues
```

Out[17]:

```
(Intercept    12.786341
CC            4.268006
Doors         6.752236
dtype: float64,
Intercept    1.580945e-35
CC           2.101878e-05
Doors        2.109558e-11
dtype: float64)
```

In [18]:

```

# Model Validation Techniques
# Two Techniques: 1. Collinearity Check & 2. Residual Analysis
# 1) Collinearity Problem Check
# Calculate VIF = 1/(1-Rsquare) for all independent variables

rsq_age=smf.ols('Age~KM+HP+CC+Doors+Gears+QT+Weight',data=toyo4).fit().rsquared
vif_age=1/(1-rsq_age)

rsq_KM=smf.ols('KM~Age+HP+CC+Doors+Gears+QT+Weight',data=toyo4).fit().rsquared
vif_KM=1/(1-rsq_KM)

rsq_HP=smf.ols('HP~Age+KM+CC+Doors+Gears+QT+Weight',data=toyo4).fit().rsquared
vif_HP=1/(1-rsq_HP)

rsq_CC=smf.ols('CC~Age+KM+HP+Doors+Gears+QT+Weight',data=toyo4).fit().rsquared
vif_CC=1/(1-rsq_CC)

rsq_DR=smf.ols('Doors~Age+KM+HP+CC+Gears+QT+Weight',data=toyo4).fit().rsquared
vif_DR=1/(1-rsq_DR)

rsq_GR=smf.ols('Gears~Age+KM+HP+CC+Doors+QT+Weight',data=toyo4).fit().rsquared
vif_GR=1/(1-rsq_GR)

rsq_QT=smf.ols('QT~Age+KM+HP+CC+Doors+Gears+Weight',data=toyo4).fit().rsquared
vif_QT=1/(1-rsq_QT)

rsq_WT=smf.ols('Weight~Age+KM+HP+CC+Doors+Gears+QT',data=toyo4).fit().rsquared
vif_WT=1/(1-rsq_WT)

# Putting the values in Dataframe format
d1={'Variables':['Age','KM','HP','CC','Doors','Gears','QT','Weight'],
    'Vif':[vif_age,vif_KM,vif_HP,vif_CC,vif_DR,vif_GR,vif_QT,vif_WT]}
Vif_df=pd.DataFrame(d1)
Vif_df

```

Out[18]:

	Variables	Vif
0	Age	1.876236
1	KM	1.757178
2	HP	1.419180
3	CC	1.163470
4	Doors	1.155890
5	Gears	1.098843
6	QT	2.295375
7	Weight	2.487180

In [19]:

```
# 2) Residual Analysis
# Test for Normality of Residuals (Q-Q Plot) using residual model (model.resid)
sm.qqplot(model.resid, line='q') # 'q' - A line is fit through the quantiles # line = '4
5' - to draw the 45-degree diagonal line
plt.title("Normal Q-Q plot of residuals")
plt.show()
```



In [20]:

```
list(np.where(model.resid > 6000))
```

Out[20]:

```
[array([109, 146, 522], dtype=int64)]
```

In [21]:

```
list(np.where(model.resid < -6000))
```

Out[21]:

```
[array([220, 600, 959], dtype=int64)]
```


In [22]:

```
# Test for Homoscedasticity or Heteroscedasticity (plotting model's standardized fitted
values vs standardized residual values)

def standard_values(vals) : return (vals - vals.mean()) / vals.std() # User defined z = (x
- mu) / sigma
```

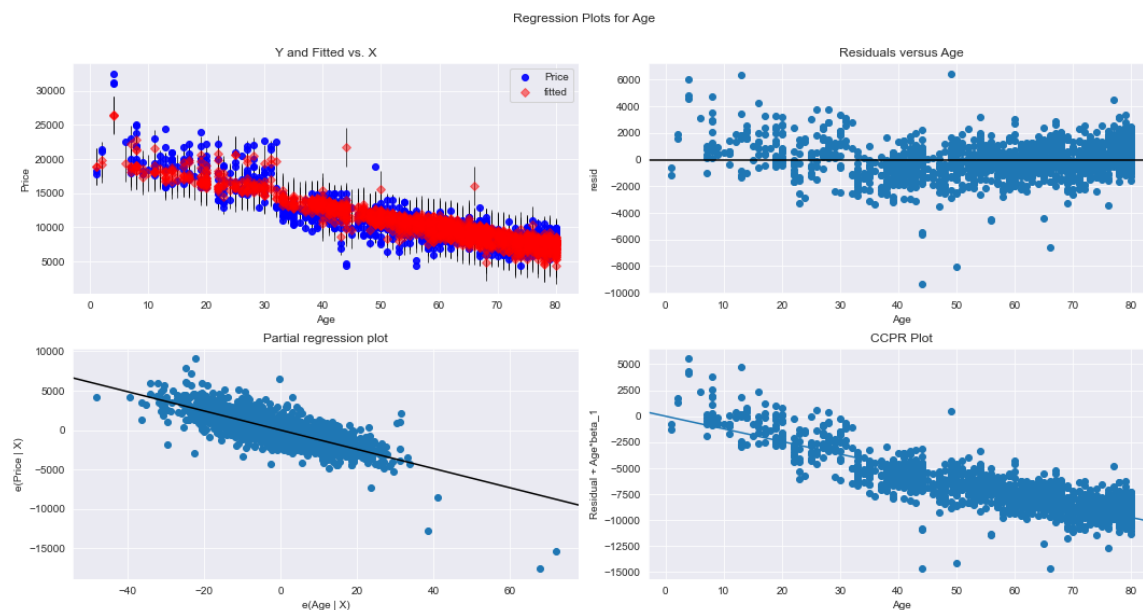
In [23]:

```
plt.scatter(standard_values(model.fittedvalues),standard_values(model.resid))
plt.title('Residual Plot')
plt.xlabel('standardized fitted values')
plt.ylabel('standardized residual values')
plt.show()
```



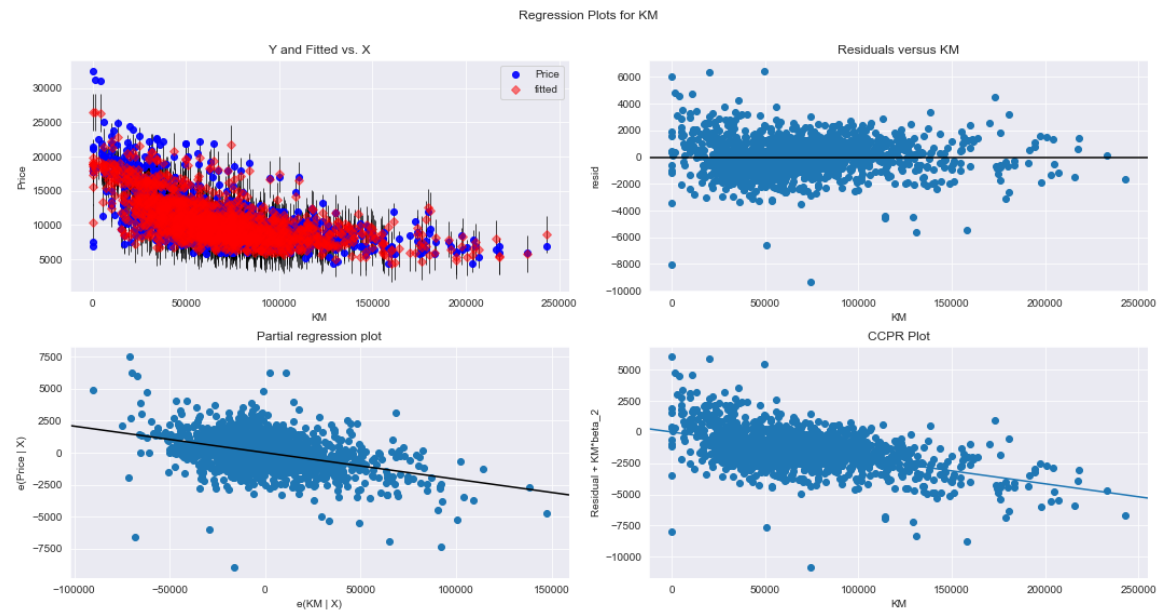
In [24]:

```
# Test for errors or Residuals Vs Regressors or independent 'x' variables or predictors
# using Residual Regression Plots code graphics.plot_regress_exog(model,'x',fig) # e
xog = x-variable & endog = y-variable
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Age',fig=fig)
plt.show()
```



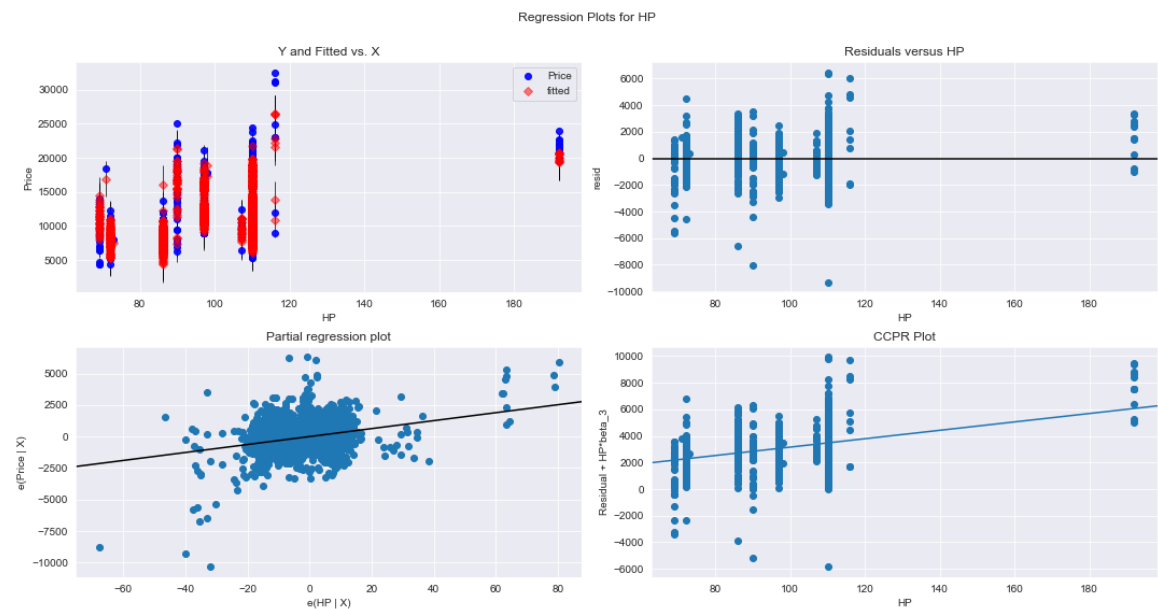
In [25]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'KM',fig=fig)
plt.show()
```



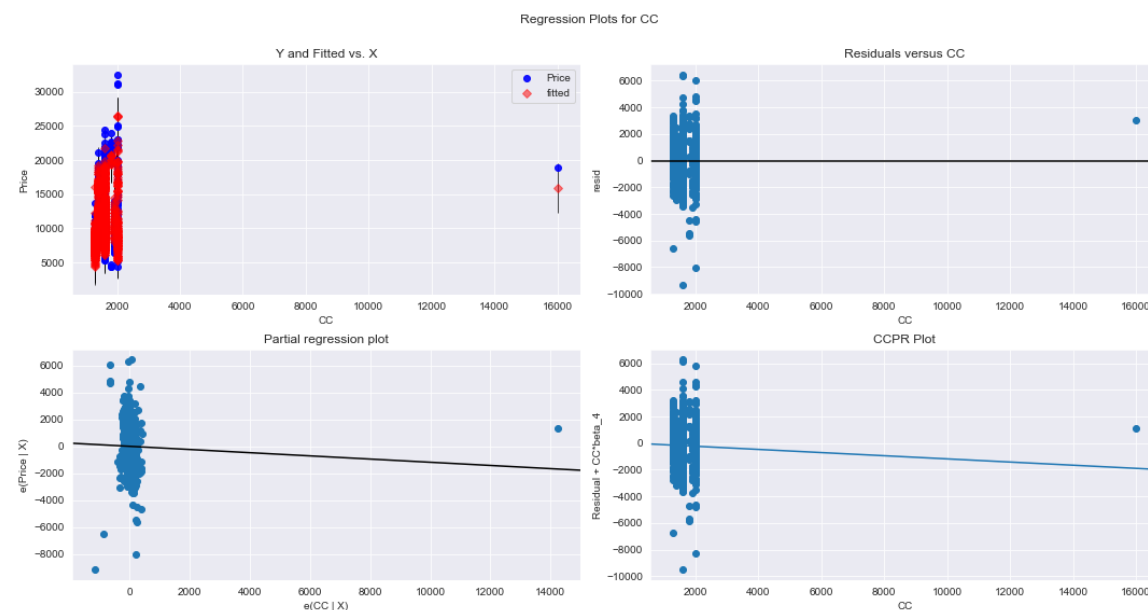
In [26]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'HP',fig=fig)
plt.show()
```



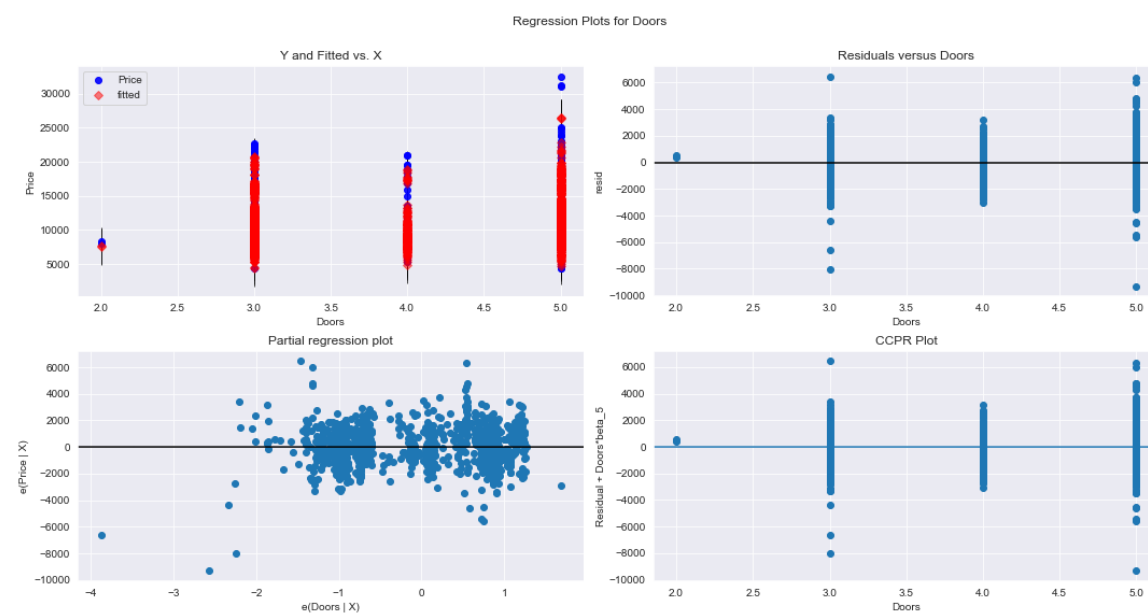
In [27]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'CC',fig=fig)
plt.show()
```



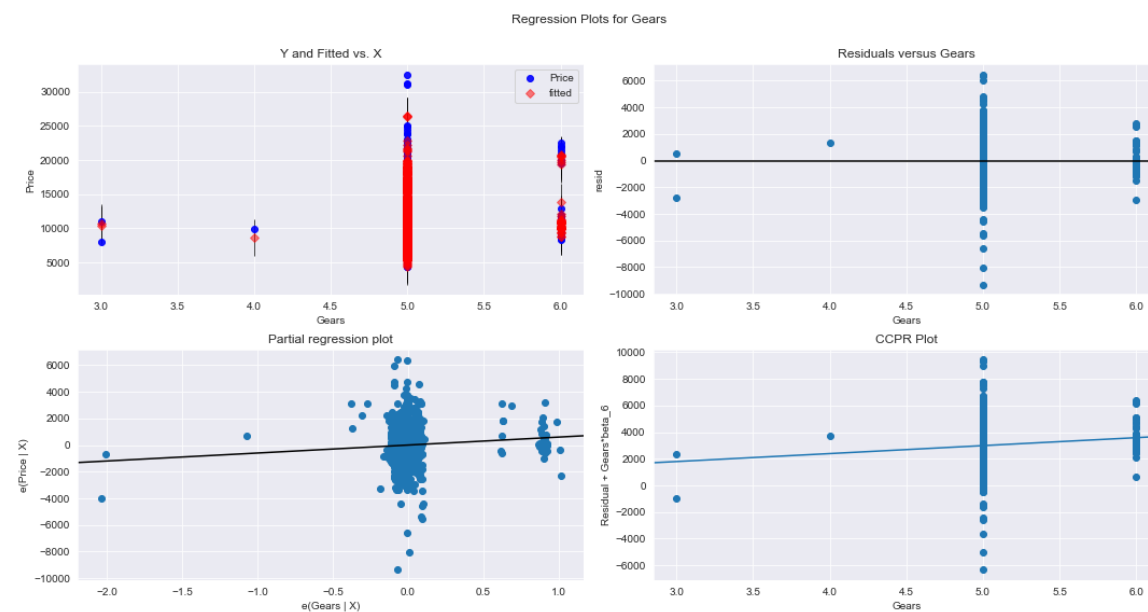
In [28]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Doors',fig=fig)
plt.show()
```



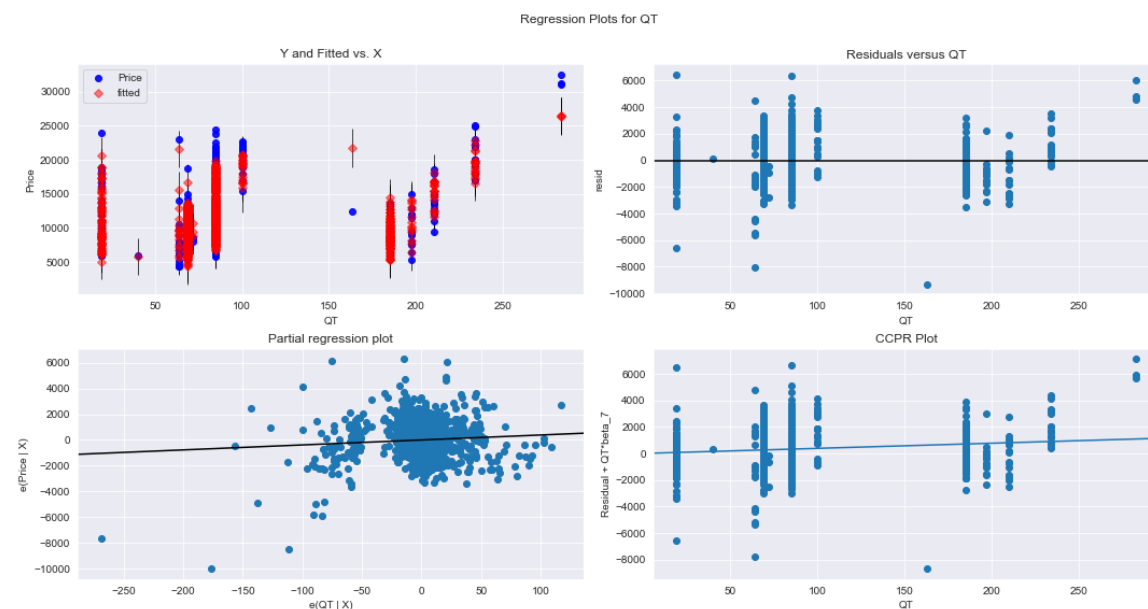
In [29]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Gears', fig=fig)
plt.show()
```



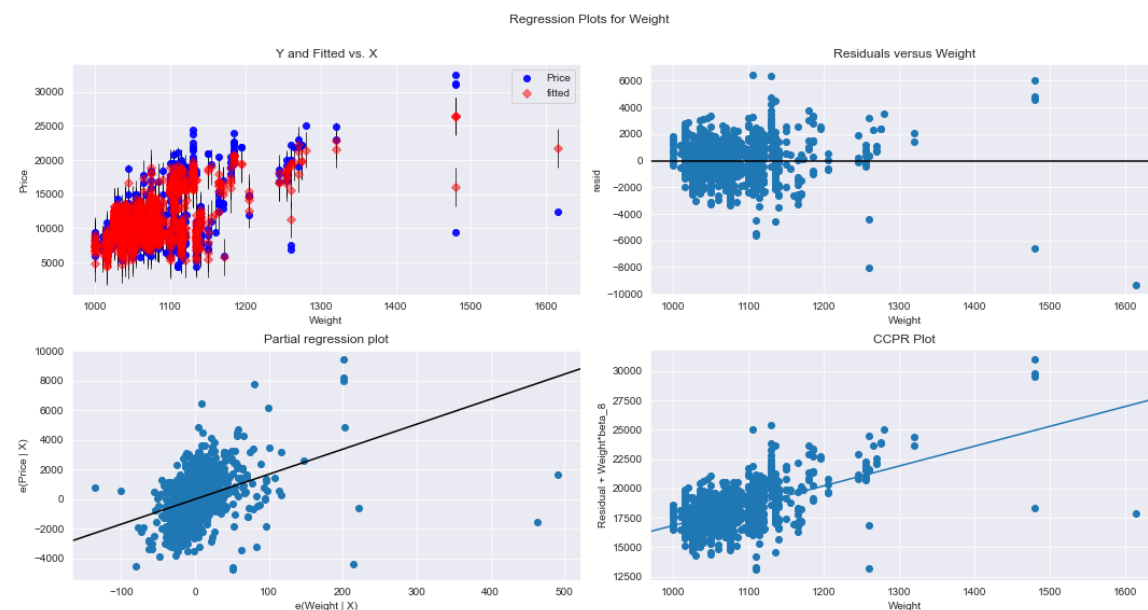
In [30]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'QT', fig=fig)
plt.show()
```



In [31]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Weight', fig=fig)
plt.show()
```



In [32]:

```
# Model Deletion Diagnostics (checking Outliers or Influencers)
# Two Techniques : 1. Cook's Distance & 2. Leverage value
# 1. Cook's Distance: If Cook's distance > 1, then it's an outlier
# Get influencers using cook's distance
(c,_)=model.get_influence().cooks_distance
c
```

Out[32]:


```
array([7.22221054e-03, 3.94547973e-03, 5.44224039e-03, ...,
       8.04110550e-07, 6.99854767e-04, 1.08408002e-02])
```

In [33]:

```
# Plot the influencers using the stem plot
fig=plt.figure(figsize=(20,7))
plt.stem(np.arange(len(toyo4)),np.round(c,3))
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance')
plt.show()
```

<ipython-input-33-86da29b54e34>:3: UserWarning: In Matplotlib 3.3 individual lines on a stem plot will be added as a LineCollection instead of individual lines. This significantly improves the performance of a stem plot. To remove this warning and switch to the new behaviour, set the "use_line_collection" keyword argument to True.

```
plt.stem(np.arange(len(toyo4)),np.round(c,3))
```



In [34]:


```
# Index and value of influencer where C>0.5
np.argmax(c) , np.max(c)
```

Out[34]:

```
(80, 78.7295058224916)
```

```
# 2. Leverage Value using High Influence Points : Points beyond Leverage_cutoff value are influencers
```

```
fig=influence_plot(model,ax = ax)
```



In [36]:

```
# Leverage Cutoff Value = 3*(k+1)/n ; k = no.of features/columns & n = no. of datapoints
k=toyo4.shape[1]
n=toyo4.shape[0]
leverage_cutoff = (3*(k+1))/n
leverage_cutoff
```

Out[36]:

0.020905923344947737

In [37]:

```
toyo4[toyo4.index.isin([80])]
```

Out[37]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
80	18950	25	20019	110	16000	5	5	100	1180

In [38]:

```
# Improving the Model
# Creating a copy of data so that original dataset is not affected
toyo_new=toyo4.copy()
toyo_new
```

Out[38]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1430	7500	69	20544	86	1300	3	5	69	1025
1431	10845	72	19000	86	1300	3	5	69	1015
1432	8500	71	17016	86	1300	3	5	69	1015
1433	7250	70	16916	86	1300	3	5	69	1015
1434	6950	76	1	110	1600	5	5	19	1114

1435 rows × 9 columns

In [39]:

```
# Discard the data points which are influencers and reassign the row number (reset_index(drop=True))
toyo5=toyo_new.drop(toyo_new.index[[80]],axis=0).reset_index(drop=True)
toyo5
```

Out[39]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1429	7500	69	20544	86	1300	3	5	69	1025
1430	10845	72	19000	86	1300	3	5	69	1015
1431	8500	71	17016	86	1300	3	5	69	1015
1432	7250	70	16916	86	1300	3	5	69	1015
1433	6950	76	1	110	1600	5	5	19	1114

1434 rows × 9 columns

In [40]:

```
# Model Deletion Diagnostics and Final Model
while np.max(c)>0.5 :
    model=smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=toyo5).fit()
    (c,_)=model.get_influence().cooks_distance
    c
    np.argmax(c) , np.max(c)
    toyo5=toyo5.drop(toyo5.index[[np.argmax(c)]],axis=0).reset_index(drop=True)
    toyo5
else:
    final_model=smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=toyo5).fit()
    final_model.rsquared , final_model.aic
    print("Thus model accuracy is improved to",final_model.rsquared)
```

Thus model accuracy is improved to 0.8882395145171204

In [41]:

```

if np.max(c)>0.5:
    model=smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=toyo5).fit()
    (c,_)=model.get_influence().cooks_distance
    c
    np.argmax(c) , np.max(c)
    toyo5=toyo5.drop(toyo5.index[[np.argmax(c)]],axis=0).reset_index(drop=True)
    toyo5
elif np.max(c)<0.5:
    final_model=smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=toyo5).fit()
    final_model.rsquared , final_model.aic
    print("Thus model accuracy is improved to",final_model.rsquared)

```

Thus model accuracy is improved to 0.8882395145171204

In [42]:

```
final_model.rsquared
```

Out[42]:

0.8882395145171204

In [43]:

```
toyo5
```

Out[43]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1426	7500	69	20544	86	1300	3	5	69	1025
1427	10845	72	19000	86	1300	3	5	69	1015
1428	8500	71	17016	86	1300	3	5	69	1015
1429	7250	70	16916	86	1300	3	5	69	1015
1430	6950	76	1	110	1600	5	5	19	1114

1431 rows × 9 columns

In [44]:

```
# Model Predictions
# say New data for prediction is
new_data=pd.DataFrame({'Age':12,"KM":40000,"HP":80,"CC":1300,"Doors":4,"Gears":5,"QT":69,"Weight":1012},index=[0])
new_data
```

Out[44]:

	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	12	40000	80	1300	4	5	69	1012

In [45]:

```
# Manual Prediction of Price
final_model.predict(new_data)
```

Out[45]:

```
0    14341.570181
dtype: float64
```

In [46]:

```
# Automatic Prediction of Price with 90.02% accuracy
pred_y=final_model.predict(toyo5)
pred_y
```

Out[46]:

```
0    16345.352610
1    15886.635544
2    16328.224968
3    15996.318854
4    15883.424182
...
1426   9161.230587
1427   8536.091326
1428   8681.531063
1429   8793.668694
1430  10860.695492
Length: 1431, dtype: float64
```

In []: